

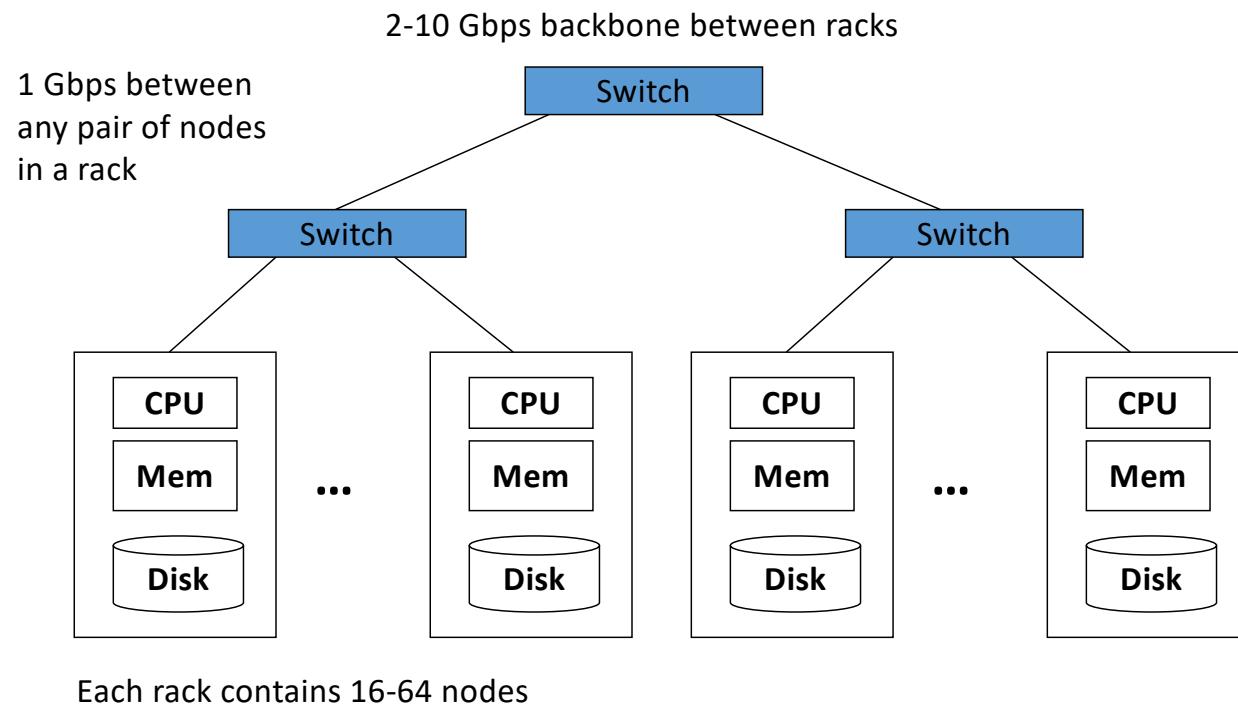
Hadoop MapReduce

Anurag Nagar, Ph.D.

Hadoop Distributed File System (HDFS)

- We learned the architecture of HDFS.
- Each of the data nodes have storage and processing capacity.
- Need a model of data processing that is **parallel**, **distributed**, **fault-tolerant**, and **efficient**.
- Want to minimize communication between nodes as it costs network bandwidth.
- Let's look at a real example on the next page

Cluster Architecture



In 2011 it was guestimated that Google had 1M machines, <http://bit.ly/Shh0RO>



Large-scale Computing

- Problems with using commodity hardware
- **Challenges:**
 - How do you distribute computation?
 - **How can we make it easy to write distributed programs?**
 - Machines fail:
 - One server may stay up 3 years (1,000 days)
 - If you have 1,000 servers, expect to loose 1/day
 - People estimated Google had ~1M machines in 2011
 - 1,000 machines fail every day!

MapReduce in HDFS

- MR is the processing engine of HDFS.
- Helps with the concept of "moving computation" rather than "moving data".
=> locality of computation
- Cluster consists of nodes, that have storage and processing power.
- We need to have multiple nodes perform computation in parallel.

MapReduce

- **Design Considerations:**

- process vast amounts of data (multi-terabyte data-sets)
- parallel processing
- large clusters (thousands of nodes) of commodity hardware
- reliable
- fault-tolerant
- should be able to increase processing power by adding more nodes
 - > "scale-out" and not "scale-up".
- sharing data or processing between nodes is bad
 - > ideally want "shared-nothing" architecture.
- want batch processing
 - > process entire dataset and not random seeks

MapReduce Basics

MR has origins in Functional Programming

- Map is a higher order function that applies a function element-wise to a list of elements.
- Map transform **lists** of **input data** elements into **lists** of **output data elements** by applying a function to each element of the list.
- Reduce (also called Fold) is a higher order function that processes a list of elements by applying a function pairwise and finally returning a scalar.
- Reduce compacts a list into a scalar by applying a function pairwise.

Functional Programming

- **Key feature: higher order functions**

- ▶ Functions that accept other functions as arguments
- ▶ **Map and Fold (Reduce)**




Figure: Illustration of *map* and *fold*.

Map Operation


- Define a function: `square x = x * x`

- Apply on a list: `>>> map square [1, 2, 3, 4, 5]`


- Get another list: `[1, 4, 9, 16, 25]`

Map function

- Takes input (k, v) and outputs (k', v')
=> Generally input k has little meaning, but we try to find a meaningful output k'
- Example: You have input file with line number as key and text as value. A map function could extract and output year as key and temperature as value.



Mapping



Mapper Process

Reduce (Fold) Operation

- Define an operator: +
- Initial value = 0
- Apply on a list: `[1,2,3,4,5]`
- Get a scalar: 15

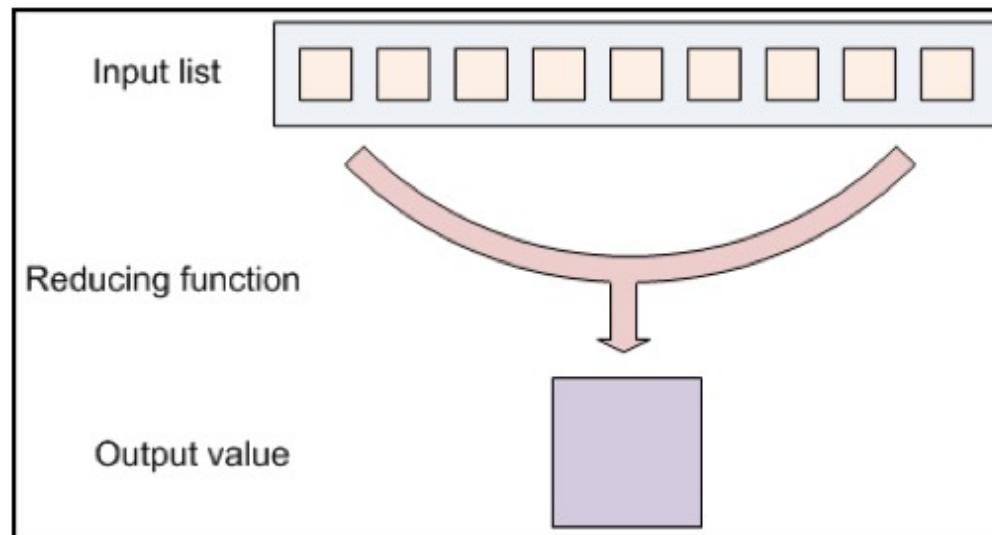
Reduce function

- Reduce function generally receives a key and a list of values.
- It compacts the list to a single (generally) value.
- For example, input key is year and value is list of temperatures. Output could be key and maximum temperature.

$(1950, [30, 70, 50, 72, 18]) \rightarrow (1950, 72)$

- A key point is that reduce is generally run on data from same key value.
=> Eg. Find average time spent by each visitor on a website
Key = userID, Value = Time spent during each visit
It makes sense to aggregate (reduce) for each key separately

Reducing



Reducer Process

MapReduce Data Structures

Key-Value Structure

- Each data element needs to have a key associated with it.
- Uniquely identifies the data item.
- Example: Log of cars passing by.
What's the key?
Could be the license plate number.

```
AAA-123  65mph, 12:00pm
ZZZ-789  50mph, 12:02pm
AAA-123  40mph, 12:05pm
CCC-456  25mph, 12:15pm
...
```

- Does it have to be unique in entire dataset? No

K-V pairs

- Key-Value (K-V) pairs are one of the basic data structures for BD.
- Please keep this in mind for future discussion also.


K-V pairs

- A mapper is presented data that contains multiple keys.
- It transforms this data in a 1-1 fashion and outputs a meaningful K-V pair.
- The reducer is presented with data containing only a single key.
- It compacts (or aggregates) the values of the key.
- How does each reducer get data from only one key?
- Someone has to do the sorting and shuffling of data from mappers to reducers.
- That's the job of the Hadoop framework




MapReduce in Hadoop

MapReduce: The Map Step



MapReduce: The Reduce Step



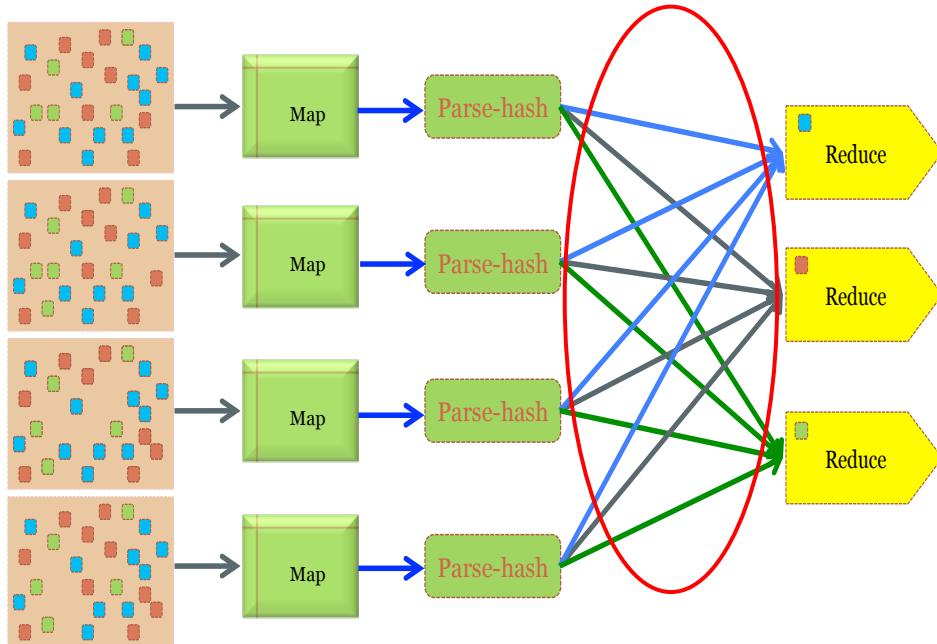
Key-Value Pairs

- Mappers and Reducers are users' code (provided functions)
- Just need to obey the Key-Value pairs interface
- **Mappers:**
 - Consume $\langle \text{key}, \text{value} \rangle$ pairs
 - Produce $\langle \text{key}, \text{value} \rangle$ pairs
- **Reducers:**
 - Consume $\langle \text{key}, \langle \text{list of values} \rangle \rangle$
 - Produce $\langle \text{key}, \text{value} \rangle$
- **Shuffling and Sorting:**
 - Hidden phase between mappers and reducers
 - Groups all similar keys from all mappers, sorts and passes them to a certain reducer in the form of $\langle \text{key}, \langle \text{list of values} \rangle \rangle$

Example 1 – Color Count

MapReduce Execution in Hadoop

- Suppose you are given a dataset where each item is keyed with a color – **Red**, **Blue**, or **Green**
- Aim is to compute the count of each colors.



Dataset is divided into 4 blocks.

The map-reduce job consists of 4 map tasks and 3 reduce tasks


Map task takes each data item and applies a transformation to it. Could be as simple as output (key, 1) e.g. (Red, 1)

Reduce task needs to get data of a single key.

Framework does the sorting and shuffling

Color Count Example

Job: Count the number of each color in a data set



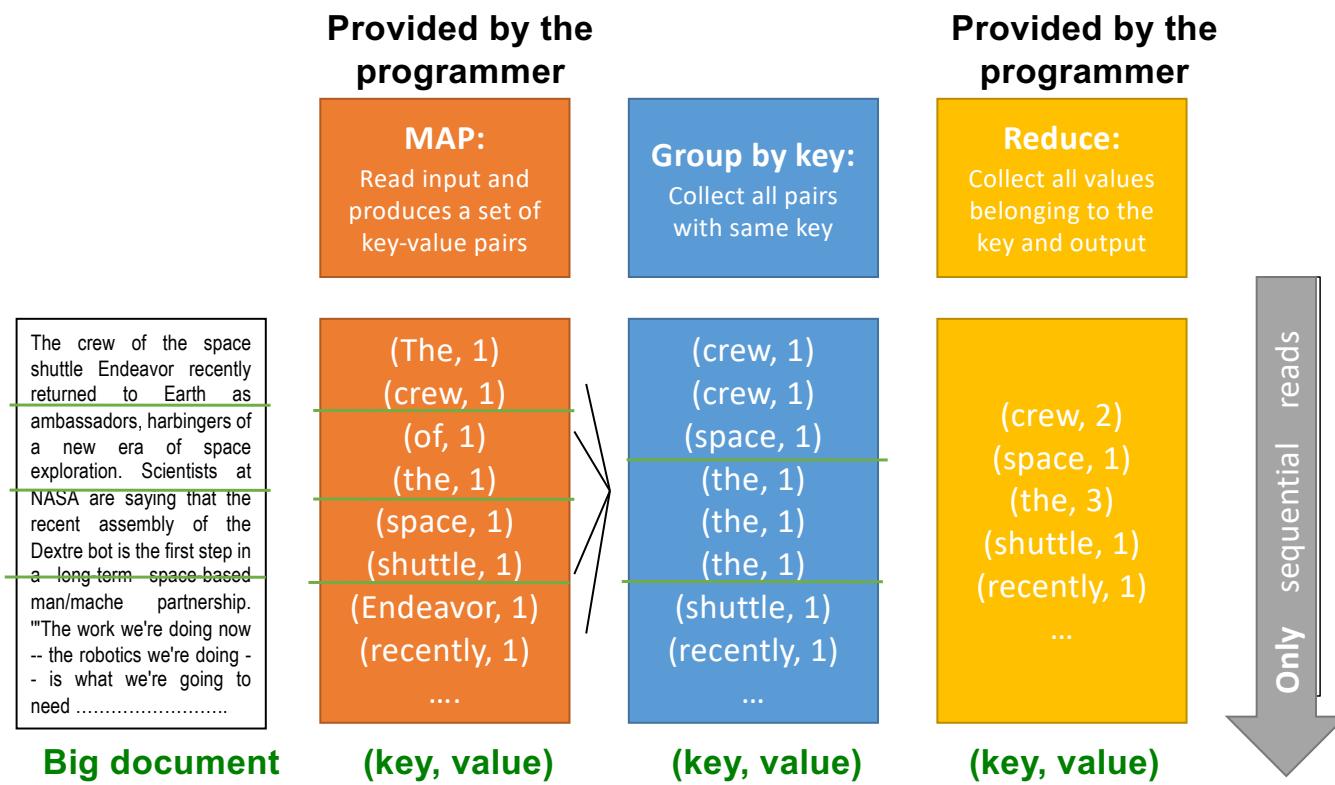
Example 2 – Word Count

Programming Model: MapReduce

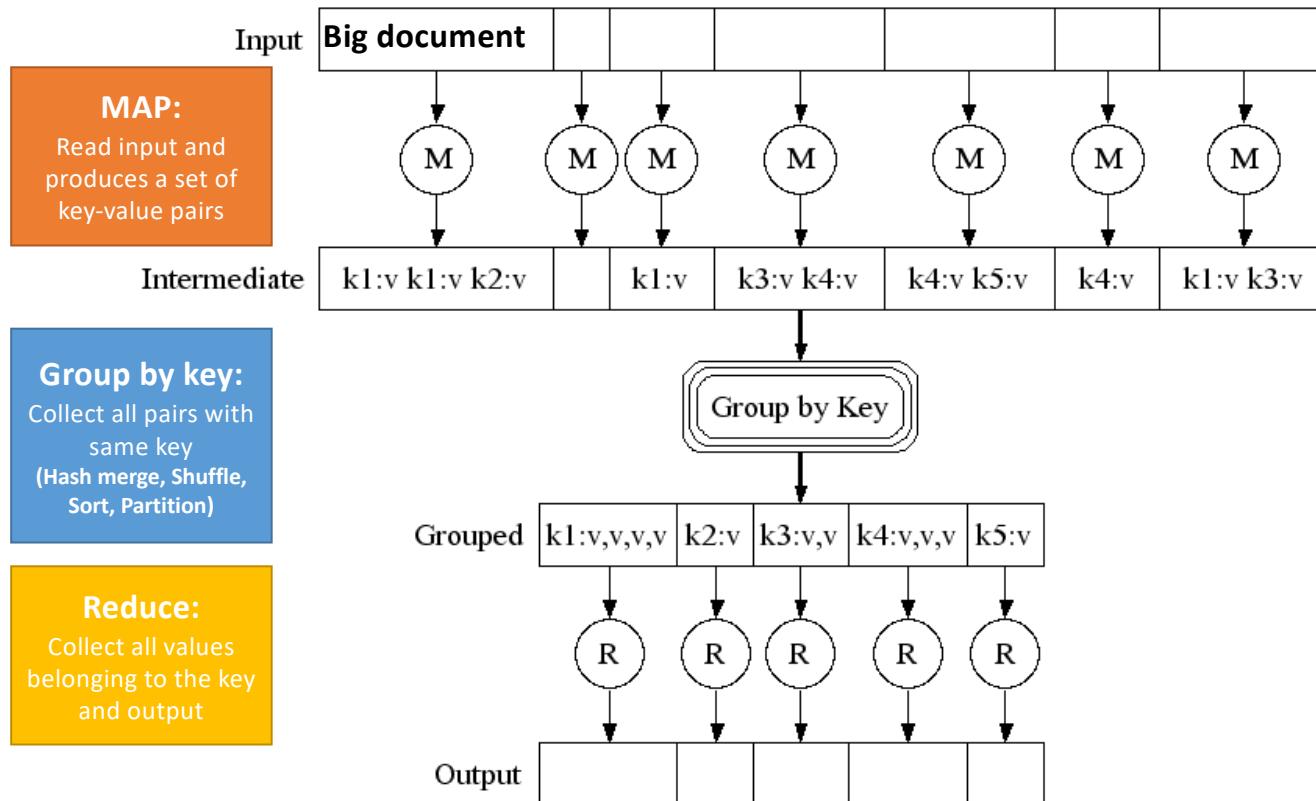
Warm-up task:

- We have a huge text document
- Count the number of times each distinct word appears in the file
- **Sample application:**
 - Analyze web server logs to find popular URLs

MapReduce: Word Counting



Map-Reduce: A diagram



Word Count Using MapReduce

```
map(key, value):  
    // key: document name; value: text of the document  
    for each word w in value:  
        emit(w, 1)  
  
reduce(key, values):  
    // key: a word; value: an iterator over counts  
    result = 0  
    for each count v in values:  
        result += v  
    emit(key, result)
```

Map-Reduce: Environment

Map-Reduce environment takes care of:

- **Partitioning** the input data (input splits)
- **Scheduling** the program's execution across a set of machines
- Performing the **group by key** step
- Handling machine **failures**
- Managing required inter-machine **communication**

Map-Reduce

- Programmer specifies:
 - Map and Reduce and input files
- **Workflow:**
 - Read inputs as a set of key-value-pairs
 - **Map** transforms input kv-pairs into a new set of $k'v'$ -pairs
 - Sorts & Shuffles the $k'v'$ -pairs to output nodes
 - All $k'v'$ -pairs with a given k' are sent to the same **reduce**
 - **Reduce** processes all $k'v'$ -pairs grouped by key into new $k''v''$ -pairs
 - Write the resulting pairs to files
- All phases are distributed with many tasks doing the work

