

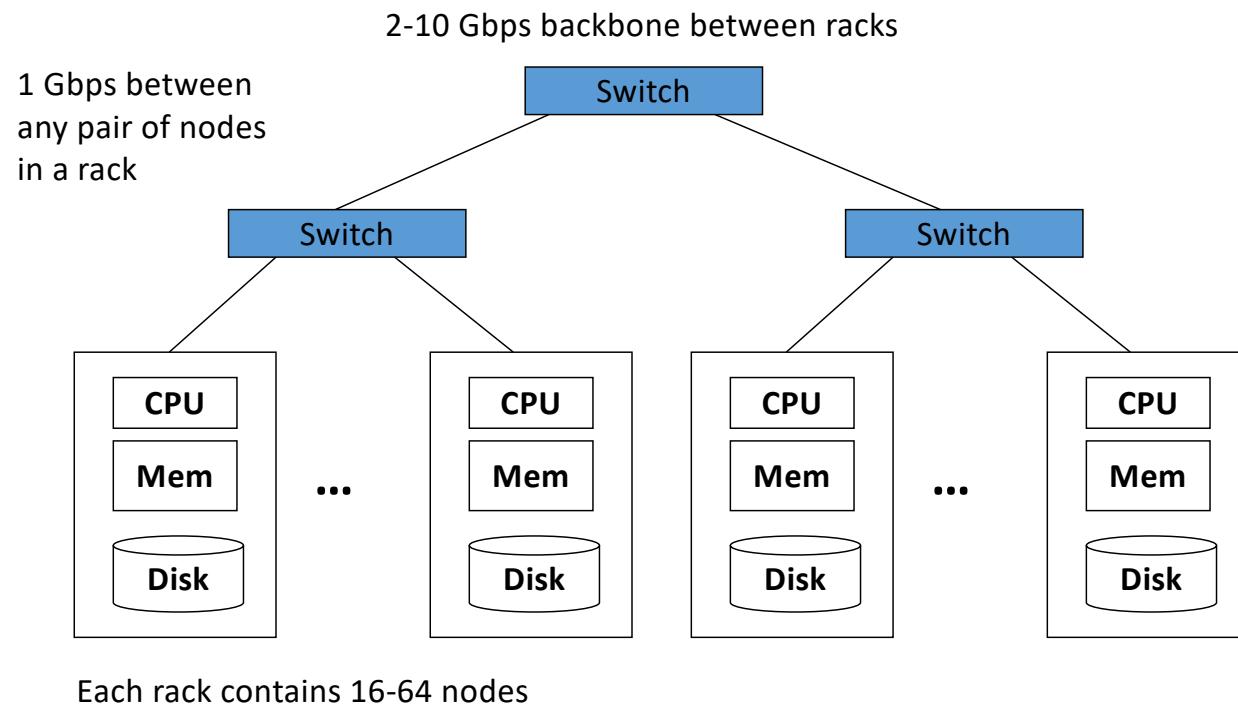
# Hadoop MapReduce

Anurag Nagar, Ph.D.

# Hadoop Distributed File System (HDFS)

- HDFS has multi-node architecture.
- Each of the data nodes have storage and processing capacity.
- Need a model of data processing that is **parallel**, **distributed**, **fault-tolerant**, and **efficient**.
- Want to minimize communication between nodes as it costs network bandwidth.
- Let's look at a real example on the next page

# Cluster Architecture



In 2011 it was guestimated that Google had 1M machines, <http://bit.ly/Shh0RO>



# Large-scale Computing

- Problems with using commodity hardware
- **Challenges:**
  - How do you distribute computation?
  - **How can we make it easy to write distributed programs?**
  - Machines fail:
    - One server may stay up 3 years (1,000 days)
    - If you have 1,000 servers, expect to loose 1/day
    - People estimated Google had ~1M machines in 2011
      - 1,000 machines fail every day!

# MapReduce in HDFS

- MR is the processing engine of HDFS.
- Helps with the concept of "moving computation" rather than "moving data".  
=> locality of computation
- Cluster consists of nodes, that have storage and processing power.
- We need to have multiple nodes perform computation in parallel.

# MapReduce

- **Design Considerations:**

- process vast amounts of data (multi-terabyte data-sets)
- parallel processing
- large clusters (thousands of nodes) of commodity hardware
- reliable
- fault-tolerant
- should be able to increase processing power by adding more nodes
  - > "scale-out" and not "scale-up".
- sharing data or processing between nodes is bad
  - > ideally want "shared-nothing" architecture.
- want batch processing
  - > process entire dataset and not random seeks

# MapReduce Basics

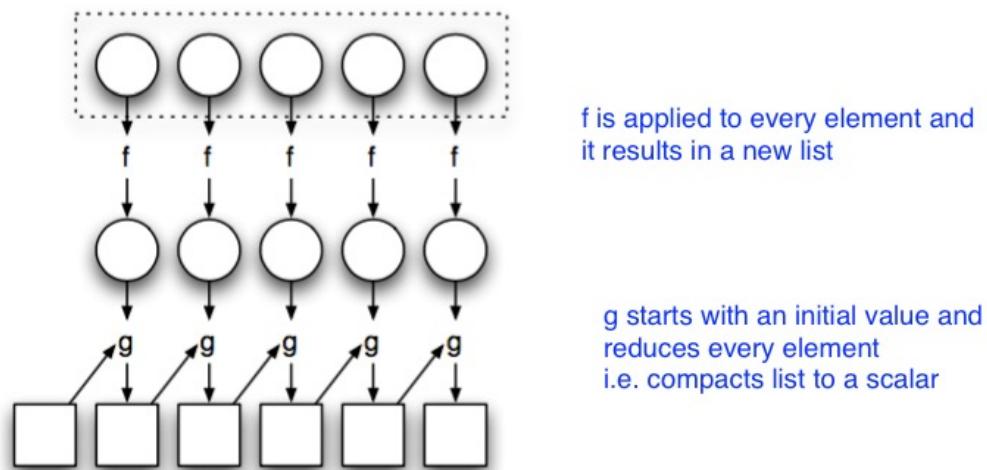
# MR has origins in Functional Programming

- Map is a higher order function that applies a function element-wise to a list of elements.
- Map transform **lists** of **input data** elements into **lists** of **output data elements** by applying a function to each element of the list.
- Reduce (also called Fold) is a higher order function that processes a list of elements by applying a function pairwise and finally returning a scalar.
- Reduce compacts a list into a scalar by applying a function pairwise.

# Functional Programming

- **Key feature: higher order functions**

- ▶ Functions that accept other functions as arguments
- ▶ **Map and Fold (Reduce)**



**Figure:** Illustration of *map* and *fold*.

# Map Operation

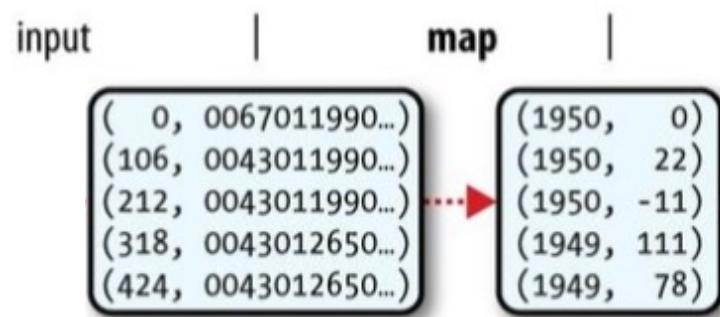
- Define a function: `square x = x * x`

- Apply on a list: `>>> map square [1, 2, 3, 4, 5]`

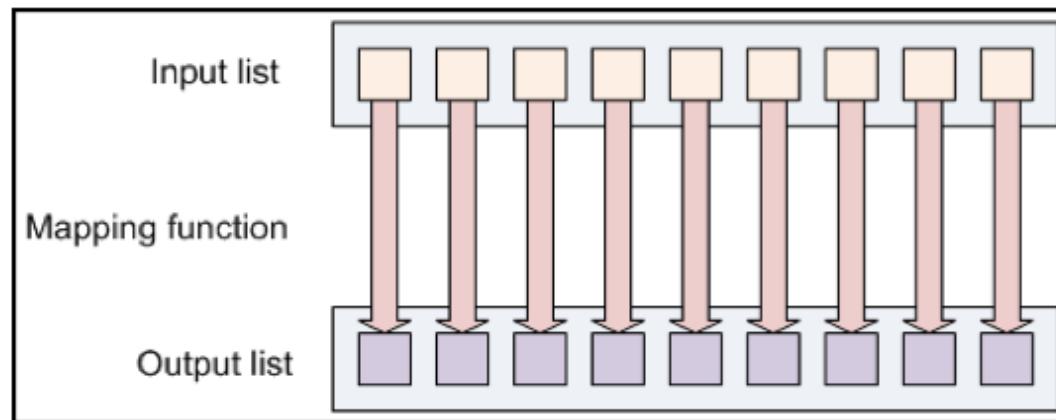
- Get another list: `[1, 4, 9, 16, 25]`

# Map function

- Takes input  $(k, v)$  and outputs  $(k', v')$   
=> Generally input  $k$  has little meaning, but we try to find a meaningful output  $k'$
- Example: You have input file with line number as key and text as value. A map function could extract and output year as key and temperature as value.



# Mapping



Mapper Process

# Reduce (Fold) Operation

- Define an operator: +
- Initial value = 0
- Apply on a list: `[1,2,3,4,5]`
- Get a scalar: 15

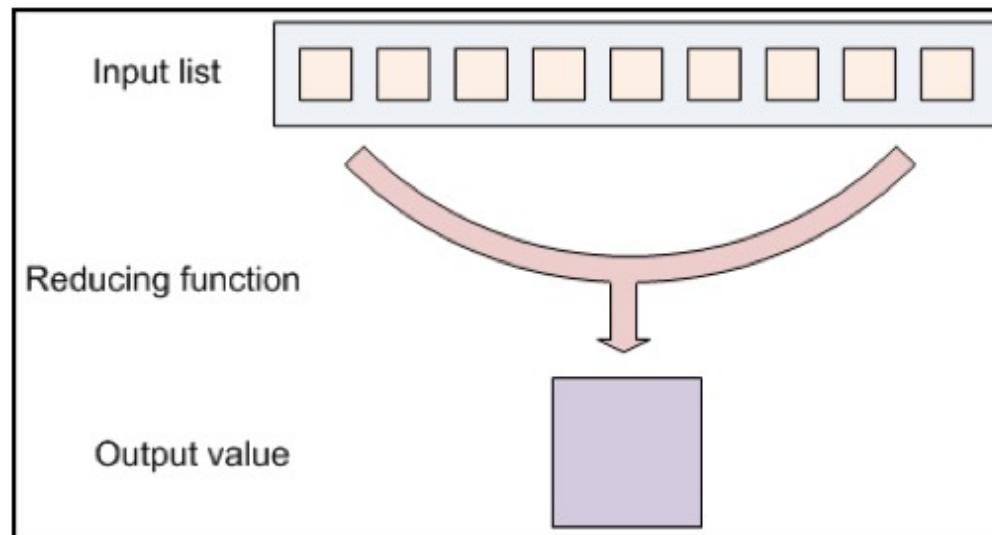
## Reduce function

- Reduce function generally receives a key and a list of values.
- It compacts the list to a single (generally) value.
- For example, input key is year and value is list of temperatures. Output could be key and maximum temperature.

$(1950, [30, 70, 50, 72, 18]) \rightarrow (1950, 72)$

- A key point is that reduce is generally run on data from same key value.  
=> Eg. Find average time spent by each visitor on a website  
Key = userID, Value = Time spent during each visit  
It makes sense to aggregate (reduce) for each key separately

# Reducing



Reducer Process

# MapReduce Data Structures

# Key-Value Structure

- Each data element needs to have a key associated with it.
- Uniquely identifies the data item.
- Example: Log of cars passing by.  
What's the key?  
Could be the license plate number.

```
AAA-123  65mph, 12:00pm
ZZZ-789  50mph, 12:02pm
AAA-123  40mph, 12:05pm
CCC-456  25mph, 12:15pm
...
```

- Does it have to be unique in entire dataset? No

## K-V pairs

- Key-Value (K-V) pairs are one of the basic data structures for BD.
- Please keep this in mind for future discussion also.

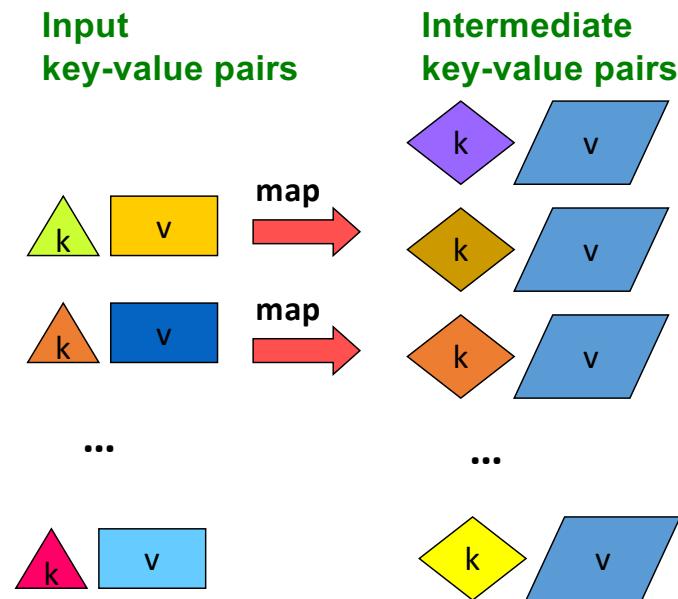
## K-V pairs

- A mapper is presented data that contains multiple keys.
- It transforms this data in a 1-1 fashion and outputs a meaningful K-V pair.
- The reducer is presented with data containing only a single key.
- It compacts (or aggregates) the values of the key.
- How does each reducer get data from only one key?
- Someone has to do the sorting and shuffling of data from mappers to reducers.
- That's the job of the Hadoop framework

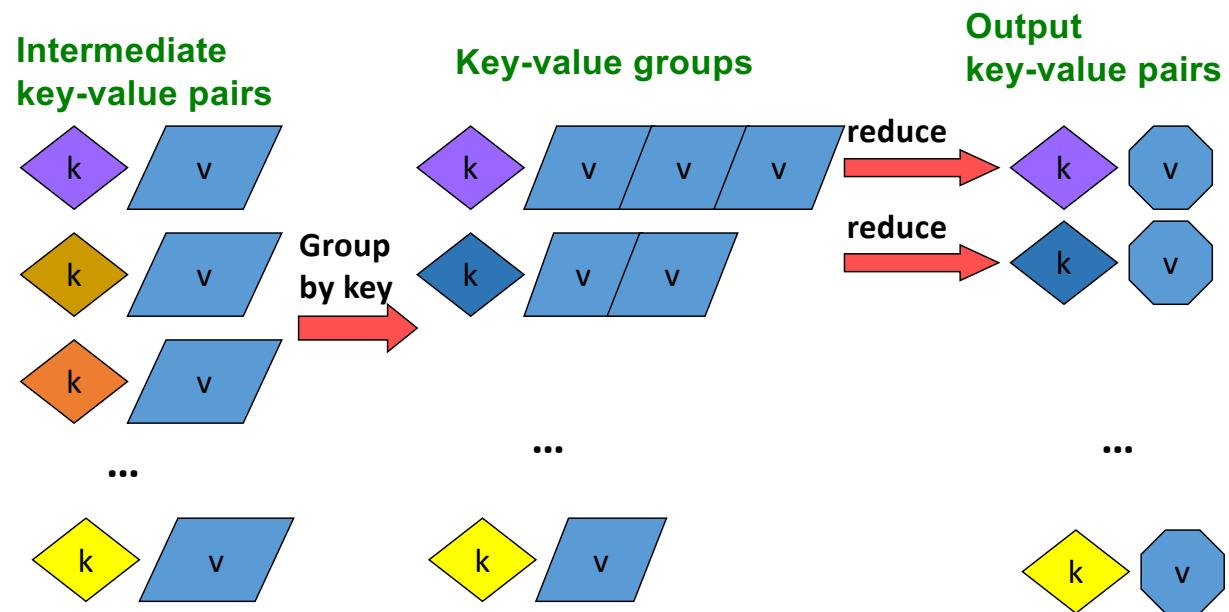


# MapReduce in Hadoop

# MapReduce: The Map Step



# MapReduce: The Reduce Step



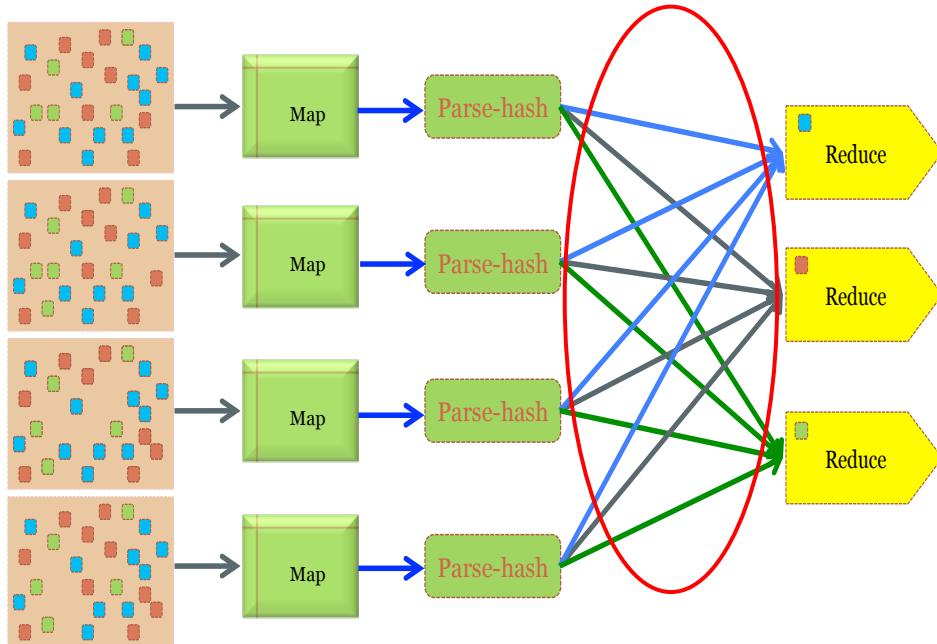
# Key-Value Pairs

- Mappers and Reducers are users' code (provided functions)
- Just need to obey the Key-Value pairs interface
- **Mappers:**
  - Consume  $\langle \text{key}, \text{value} \rangle$  pairs
  - Produce  $\langle \text{key}, \text{value} \rangle$  pairs
- **Reducers:**
  - Consume  $\langle \text{key}, \langle \text{list of values} \rangle \rangle$
  - Produce  $\langle \text{key}, \text{value} \rangle$
- **Shuffling and Sorting:**
  - Hidden phase between mappers and reducers
  - Groups all similar keys from all mappers, sorts and passes them to a certain reducer in the form of  $\langle \text{key}, \langle \text{list of values} \rangle \rangle$

# Example 1 – Color Count

# MapReduce Execution in Hadoop

- Suppose you are given a dataset where each item is keyed with a color – **Red**, **Blue**, or **Green**
- Aim is to compute the count of each colors.



*Dataset is divided into 4 blocks.*

*The map-reduce job consists of 4 map tasks and 3 reduce tasks*

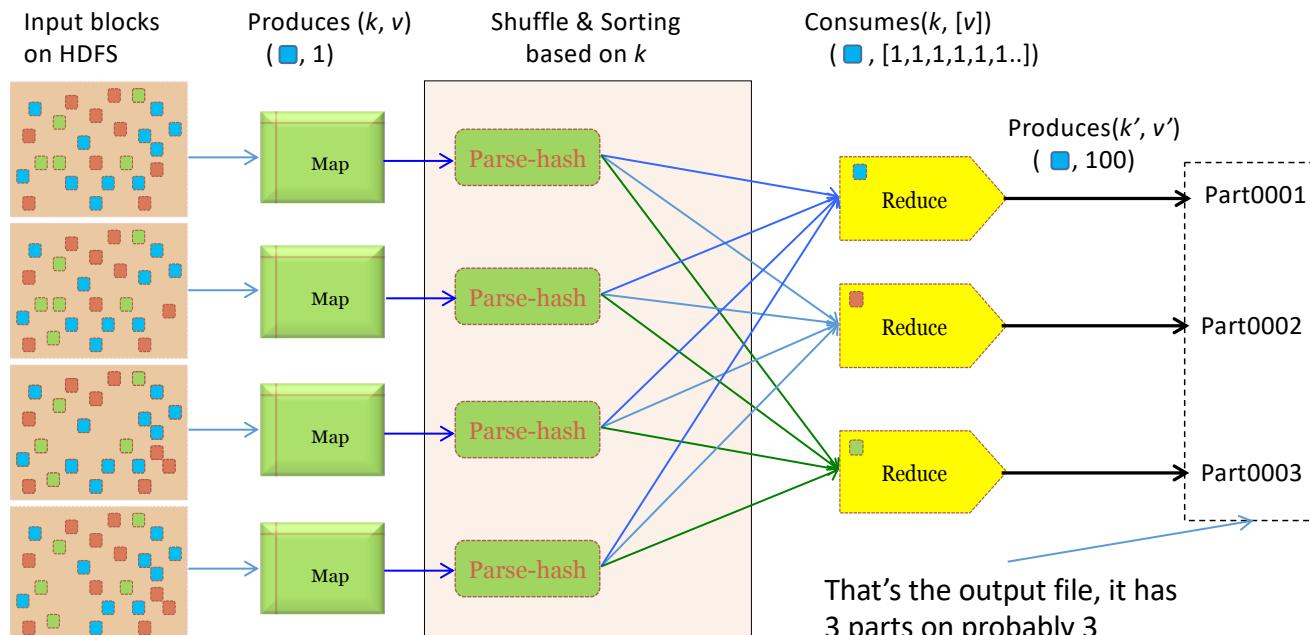
*Map task takes each data item and applies a transformation to it. Could be as simple as output (key, 1) e.g. (Red, 1)*

*Reduce task needs to get data of a single key.*

*Framework does the sorting and shuffling*

# Color Count Example

**Job: Count the number of each color in a data set**



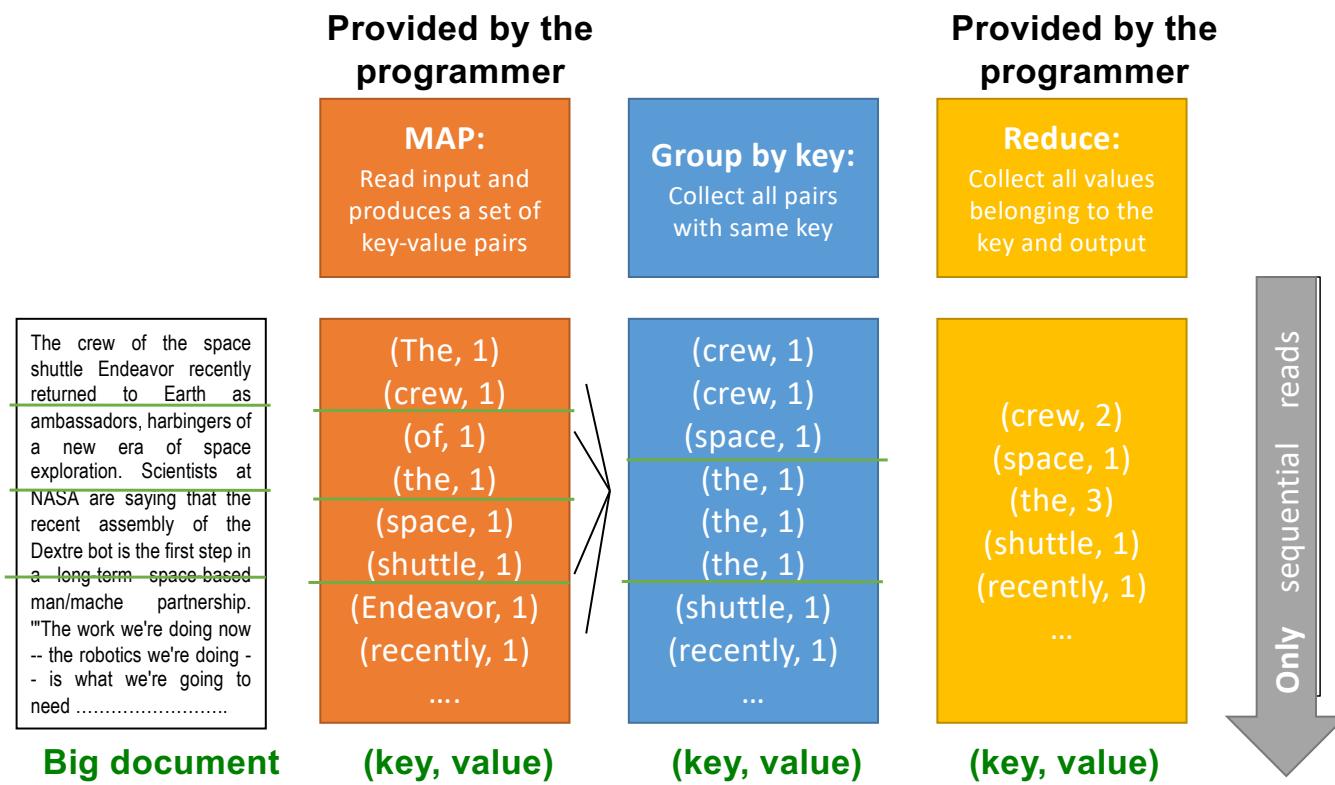
## Example 2 – Word Count

# Programming Model: MapReduce

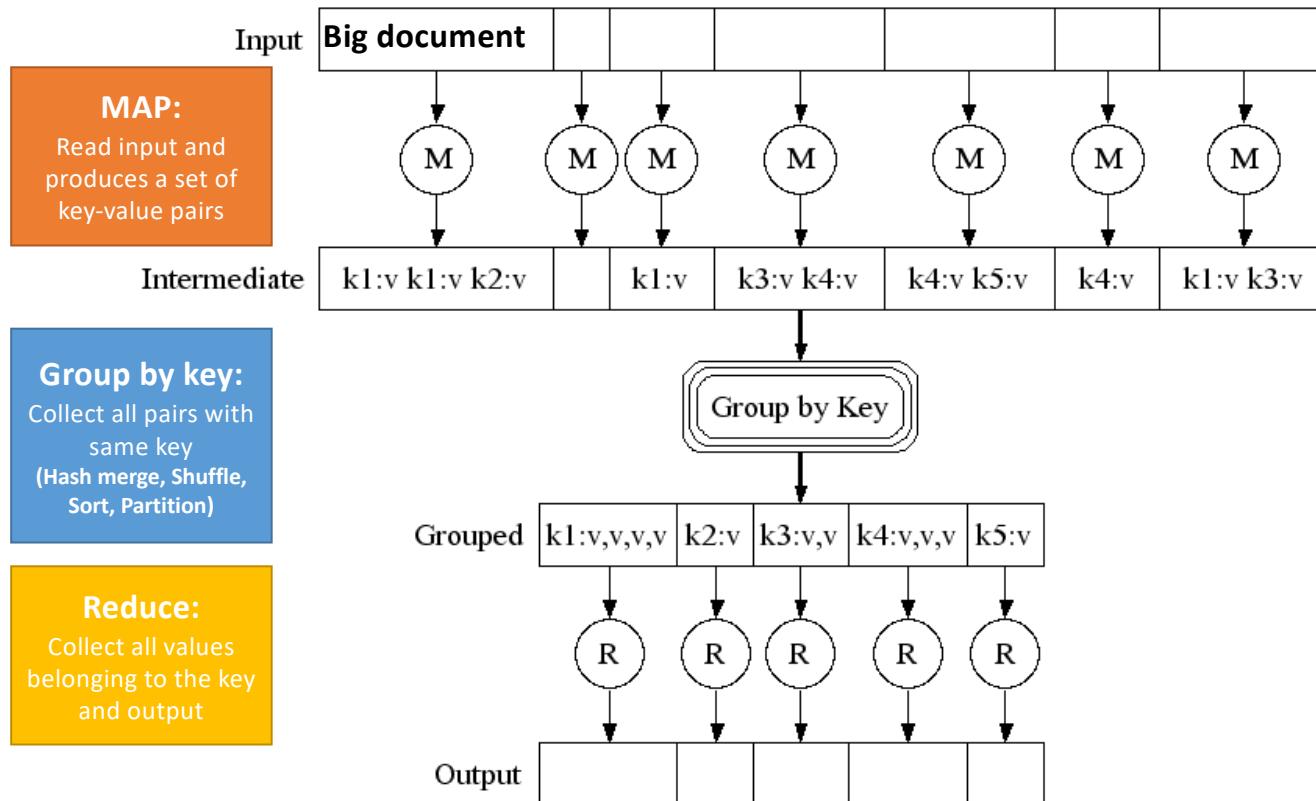
## **Warm-up task:**

- We have a huge text document
- Count the number of times each distinct word appears in the file
- **Sample application:**
  - Analyze web server logs to find popular URLs

# MapReduce: Word Counting



# Map-Reduce: A diagram



# Word Count Using MapReduce

```
map(key, value):  
    // key: document name; value: text of the document  
    for each word w in value:  
        emit(w, 1)  
  
reduce(key, values):  
    // key: a word; value: an iterator over counts  
    result = 0  
    for each count v in values:  
        result += v  
    emit(key, result)
```

# Map-Reduce: Environment

## **Map-Reduce environment takes care of:**

- **Partitioning** the input data (input splits)
- **Scheduling** the program's execution across a set of machines
- Performing the **group by key** step
- Handling machine **failures**
- Managing required inter-machine **communication**

# Map-Reduce

- Programmer specifies:
  - Map and Reduce and input files
- **Workflow:**
  - Read inputs as a set of key-value-pairs
  - **Map** transforms input kv-pairs into a new set of  $k'v'$ -pairs
  - Sorts & Shuffles the  $k'v'$ -pairs to output nodes
  - All  $k'v'$ -pairs with a given  $k'$  are sent to the same **reduce**
  - **Reduce** processes all  $k'v'$ -pairs grouped by key into new  $k''v''$ -pairs
  - Write the resulting pairs to files
- All phases are distributed with many tasks doing the work

